

Quick Links



[Online Demo](#)



[Download](#)



[FAQ](#)



[Latest Docs](#)



[Ask a
Question](#)



[Report a Bug](#)



[Screenshots](#)

This page in other versions: [Latest \(4.19\)](#) | [4.18](#) | [4.17](#) | [3.6](#) | [2.1](#) | [1.6](#) | [Development](#)

[pgAdmin 4 4.19 documentation](#) » [Backup and Restore](#) »

[previous](#) | [next](#) | [index](#)

Warning: This documentation is for a pre-release version of pgAdmin 4

Contents

- [Getting Started](#)
- [Connecting To A Server](#)
- [Managing Cluster Objects](#)
- [Managing Database Objects](#)
- [Creating or Modifying a Table](#)
- [Management Basics](#)
- [Backup and Restore](#)
 - [Backup Dialog](#)
 - [Backup Globals Dialog](#)
 - [Backup Server Dialog](#)
 - [Restore Dialog](#)
- [Developer Tools](#)
- [pgAgent](#)
- [pgAdmin Project Contributions](#)
- [Release Notes](#)
- [Licence](#)

Backup Dialog

Using the *pg_dump* utility, *pgAdmin* provides an easy way to create a backup in a plain-text or archived format. You can then use a client application (like *psql* or the *Query Tool*) to restore a plain-text backup file, or use the Postgres *pg_restore* utility to restore an archived backup. The *pg_dump* utility must have read access to all database objects that you want to back up.

You can backup a single table, a schema, or a complete database. Select the name of the backup source in the *pgAdmin* tree control, right click to open the context menu, and select *Backup...* to open the *Backup* dialog. The name of the object selected will appear in the dialog title bar.

Backup (Database: edb)

General

Dump options

Filename

...

Format

Custom

Compression ratio

Encoding

Select from the list

Number of jobs

Role name

Select from the list

i

?

✕ Cancel

Backup

Use the fields in the *General* tab to specify parameters for the backup:

- Enter the name of the backup file in the *Filename* field. Optionally, select the *Browser* icon (...) to the right to navigate into a directory and select a file that will contain the archive.
- Use the drop-down listbox in the *Format* field to select the format that is best suited for your application. Each format has advantages and disadvantages:
 - Select *Custom* to create a custom archive file that you can use with *pg_restore* to create a copy of a database. Custom archive file formats must be restored with *pg_restore*. This format offers the opportunity to select which database objects to restore from the backup file. *Custom* archive format is recommended for medium to large databases as it is compressed by default.
 - Select *Tar* to generate a tar archive file that you can restore with *pg_restore*. The tar format does not support compression.
 - Select *Plain* to create a plain-text script file. A plain-text script file contains SQL statements and commands that you can execute at the *psql* command line to recreate the database objects and load the table data. A plain-text backup file can be edited in a text editor, if desired, before using the *psql* program to restore database objects. *Plain* format is normally recommended for smaller databases; script dumps are not recommended for blobs. The SQL commands within the script will reconstruct the database to the last saved state of the database. A plain-text script can be used to reconstruct the database on another machine, or (with modifications) on other architectures.
 - Select *Directory* to generate a directory-format archive suitable for use with *pg_restore*. This file format creates a directory with one file for each table and blob being dumped, plus a *Table of Contents* file describing the dumped objects in a machine-readable format that *pg_restore* can read. This format is compressed by default.
- Use the *Compression Ratio* field to select a compression level for the backup. Specify a value of zero to mean use no compression; specify a maximum compression value of 9. Please note that tar archives do not support compression.
- Use the *Encoding* drop-down listbox to select the character encoding method that should be used for the archive.
- Use the *Number of Jobs* field (when applicable) to specify the number of tables that will be dumped simultaneously in a parallel backup.
- Use the dropdown listbox next to *Rolename* to specify the role that owns the backup.

Click the *Dump options* tab to continue. Use the box fields in the *Dump options* tab to provide options for *pg_dump*.

Backup (Table: test)

General **Dump options**

Sections

Pre-data	<input type="checkbox"/> No
Data	<input type="checkbox"/> No
Post-data	<input type="checkbox"/> No

i **?** **Cancel** **Backup**

- Move switches in the **Sections** field box to select a portion of the object that will be backed up.
 - Move the switch next to *Pre-data* to the *Yes* position to include all data definition items not included in the data or post-data item lists.
 - Move the switch next to *Data* to the *Yes* position to backup actual table data, large-object contents, and sequence values.
 - Move the switch next to *Post-data* to the *Yes* position to include definitions of indexes, triggers, rules, and constraints other than validated check constraints.

Backup (Table: test)

General **Dump options**

Type of objects

Only data	<input type="checkbox"/> No
Only schema	<input type="checkbox"/> No
Blobs	<input checked="" type="checkbox"/> Yes

i **?** **Cancel** **Backup**

- Move switches in the **Type of objects** field box to specify details about the type of objects that will be backed up.
 - Move the switch next to *Only data* to the *Yes* position to limit the back up to data.
 - Move the switch next to *Only schema* to limit the back up to schema-level database objects.
 - Move the switch next to *Blobs* to the *No* position to exclude large objects in the backup.

Backup (Table: test)

General Dump options

Do not save

Owner

Privilege

Tablespace

Unlogged table data

i **?**

- Move switches in the **Do not save** field box to select the objects that will not be included in the backup.

- Move the switch next to *Owner* to the *Yes* position to exclude commands that set object ownership.
- Move the switch next to *Privilege* to the *Yes* position to exclude commands that create access privileges.
- Move the switch next to *Tablespace* to the *Yes* position to exclude tablespaces.
- Move the switch next to *Unlogged table data* to the *Yes* position to exclude the contents of unlogged tables.
- Move the switch next to *Comments* to the *Yes* position to exclude commands that set the comments. **Note:** This option is visible only for database server greater than or equal to 11.

Backup (Table: test)

General Dump options

Queries

Use Column Inserts

Use Insert Commands

Include CREATE DATABASE statement

Include DROP DATABASE statement

i **?**

- Move switches in the **Queries** field box to specify the type of statements that should be included in the backup.

- Move the switch next to *Use Column Inserts* to the *Yes* position to dump the data in the form of INSERT statements and include explicit column names. Please note: this may make restoration from backup slow.
- Move the switch next to *Use Insert commands* to the *Yes* position to dump the data in the form of INSERT statements rather than using a COPY command. Please note: this may make restoration from backup slow.
- Move the switch next to *Include CREATE DATABASE statement* to the *Yes* position to include a command in the backup that creates a new database when restoring the backup.
- Move the switch next to *Include DROP DATABASE statement* to the *Yes* position to include a command in the backup that will drop any existing database object with the same name before recreating the object during a backup.
- Move the switch next to *Load Via Partition Root* to the *Yes* position, so when dumping a COPY or INSERT statement for a partitioned table, target the root of the partitioning hierarchy which contains it rather than the partition itself. **Note:** This option is visible only for database server greater than or equal to 11.

The screenshot shows a dialog box titled "Backup (Table: test)". It has two tabs: "General" and "Dump options", with "Dump options" currently selected. Under the "Disable" heading, there are two settings: "Trigger" and "\$ quoting". Each has a switch control currently set to "No". At the bottom of the dialog, there are three buttons: an information icon (i), a question mark icon (?), and a "Cancel" button. To the right of the "Cancel" button is a "Backup" button with a document icon.

- Move switches in the **Disable** field box to specify the type of statements that should be excluded from the backup.
 - Move the switch next to *Trigger* (active when creating a data-only backup) to the *Yes* position to include commands that will disable triggers on the target table while the data is being loaded.
 - Move the switch next to *\$ quoting* to the *Yes* position to enable dollar quoting within function bodies; if disabled, the function body will be quoted using SQL standard string syntax.

Backup (Table: test)

General **Dump options**

Miscellaneous

With OID(s) ☐ No

Verbose messages ☒ Yes

Force double quote on identifiers ☐ No

Use SET SESSION AUTHORIZATION ☐ No

Buttons: ? Cancel Backup

- Move switches in the **Miscellaneous** field box to specify miscellaneous backup options.
 - Move the switch next to *With OIDs* to the *Yes* position to include object identifiers as part of the table data for each table.
 - Move the switch next to *Verbose messages* to the *No* position to instruct *pg_dump* to exclude verbose messages.
 - Move the switch next to *Force double quotes on identifiers* to the *Yes* position to force the quoting of all identifiers.
 - Move the switch next to *Use SET SESSION AUTHORIZATION* to the *Yes* position to include a statement that will use a SET SESSION AUTHORIZATION command to determine object ownership (instead of an ALTER OWNER command).

When you've specified the details that will be incorporated into the *pg_dump* command:

- Click the *Backup* button to build and execute a command that builds a backup based on your selections on the *Backup* dialog.
- Click the *Cancel* button to exit without saving work.

Backing up an object on the server

Backing up an object on the server 'EDB Postgres Advanced Server 10 (localhost:5444)' from database 'edb'

Wed Feb 13 2019 14:01:46 GMT+0530 (India Standard Time)

⌚ 0.41 seconds More details... Stop Process

✓ Successfully completed.

Use the **Stop Process** button to stop the Backup process.

If the backup is successful, a popup window will confirm success. Click *Click here for details* on the popup window to launch the *Process Watcher*. The *Process Watcher* logs all the activity associated with the backup and provides additional information for troubleshooting.

Process Watcher - Backing up an object on the server

Backing up an object on the server 'EDB Postgres Advanced Server 10 (localhost:5444)' from database 'edb'...

Running command:

/opt/edb/as10/bin/pg_dump --file "/root/file" --host "localhost" --port "5444" --username "enterprisedb" --no-password --verbose --format=c --blobs "edb"

🕒

Start time: Wed Feb 13 2019 14:05:28 GMT+0530 (India Standard Time)

⛔ Stop Process

pg_dump: last built-in OID is 16383
pg_dump: reading extensions
pg_dump: identifying extension members
pg_dump: reading schemas
pg_dump: reading user-defined tables
pg_dump: reading packages variables
pg_dump: reading user-defined packages
pg_dump: reading user-defined synonyms
pg_dump: reading user-defined functions
pg_dump: reading user-defined types
pg_dump: finding typelem dependencies for user-defined types
pg_dump: reading procedural languages
pg_dump: reading user-defined aggregate functions
pg_dump: reading user-defined operators

✔

Successfully completed.

Execution time: 0.17 seconds

If the backup is unsuccessful, you can review the error messages returned by the backup command on the *Process Watcher*.