

Quick Links



[Online Demo](#)



[Download](#)



[FAQ](#)



[Latest Docs](#)



[Ask a
Question](#)



[Report a Bug](#)



[Screenshots](#)

This page in other versions: [Latest \(4.19\)](#) | [4.18](#) | [4.17](#) | [3.6](#) | [2.1](#) | [1.6](#) | [Development](#)

[pgAdmin 4 4.19 documentation](#) » [Creating or Modifying a Table](#) » [previous](#) | [next](#) | [index](#)

Warning: This documentation is for a pre-release version of pgAdmin 4

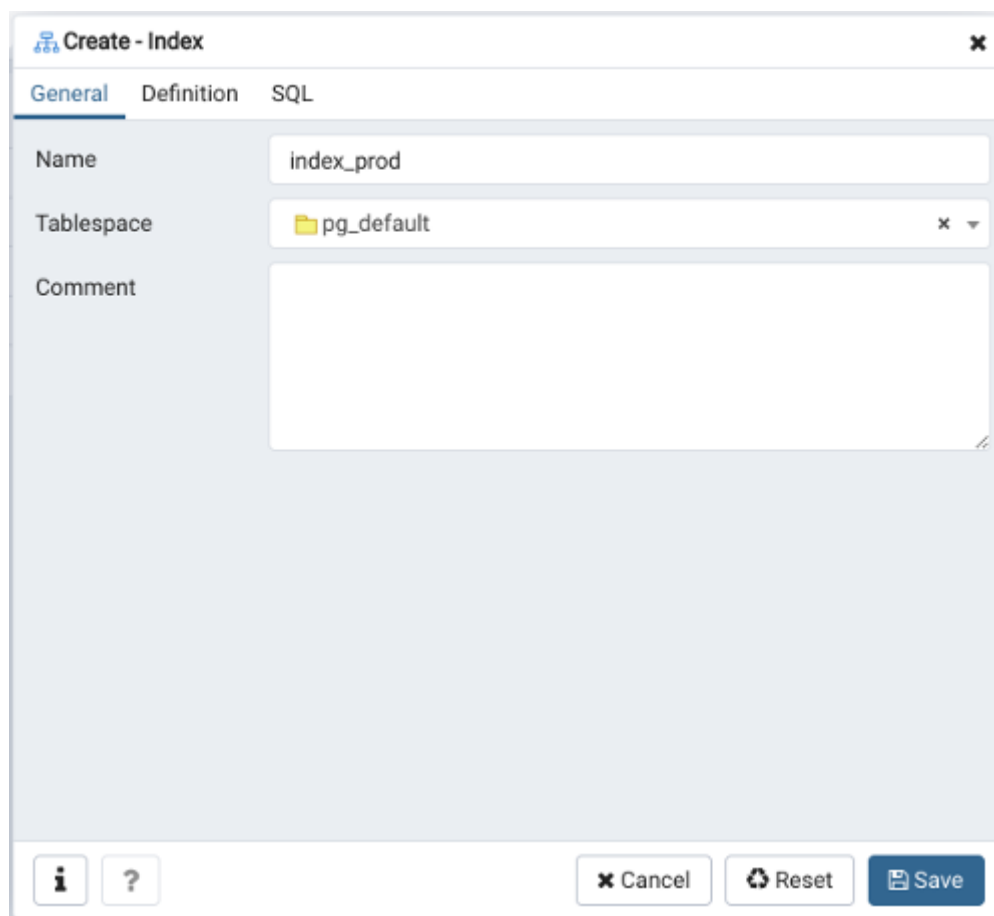
Contents

- [Index Dialog](#)
 - [Example](#)
- [Getting Started](#)
- [Connecting To A Server](#)
- [Managing Cluster Objects](#)
- [Managing Database Objects](#)
- [Creating or Modifying a Table](#)
 - [Check Dialog](#)
 - [Column Dialog](#)
 - [Compound Trigger Dialog](#)
 - [Exclusion Constraint Dialog](#)
 - [Foreign key Dialog](#)
 - [Index Dialog](#)
 - [Primary key Dialog](#)
 - [Rule Dialog](#)
 - [Table Dialog](#)
 - [Trigger Dialog](#)
 - [Unique Constraint Dialog](#)
- [Management Basics](#)
- [Backup and Restore](#)
- [Developer Tools](#)
- [pgAgent](#)
- [pgAdmin Project Contributions](#)
- [Release Notes](#)
- [Licence](#)

Index Dialog

Use the *Index* dialog to create an index on a specified table or materialized view.

The *Index* dialog organizes the development of a index through the following dialog tabs: *General* and *Definition*. The *SQL* tab displays the SQL code generated by dialog selections.



The 'Create - Index' dialog box is shown with the 'General' tab selected. It contains the following fields:

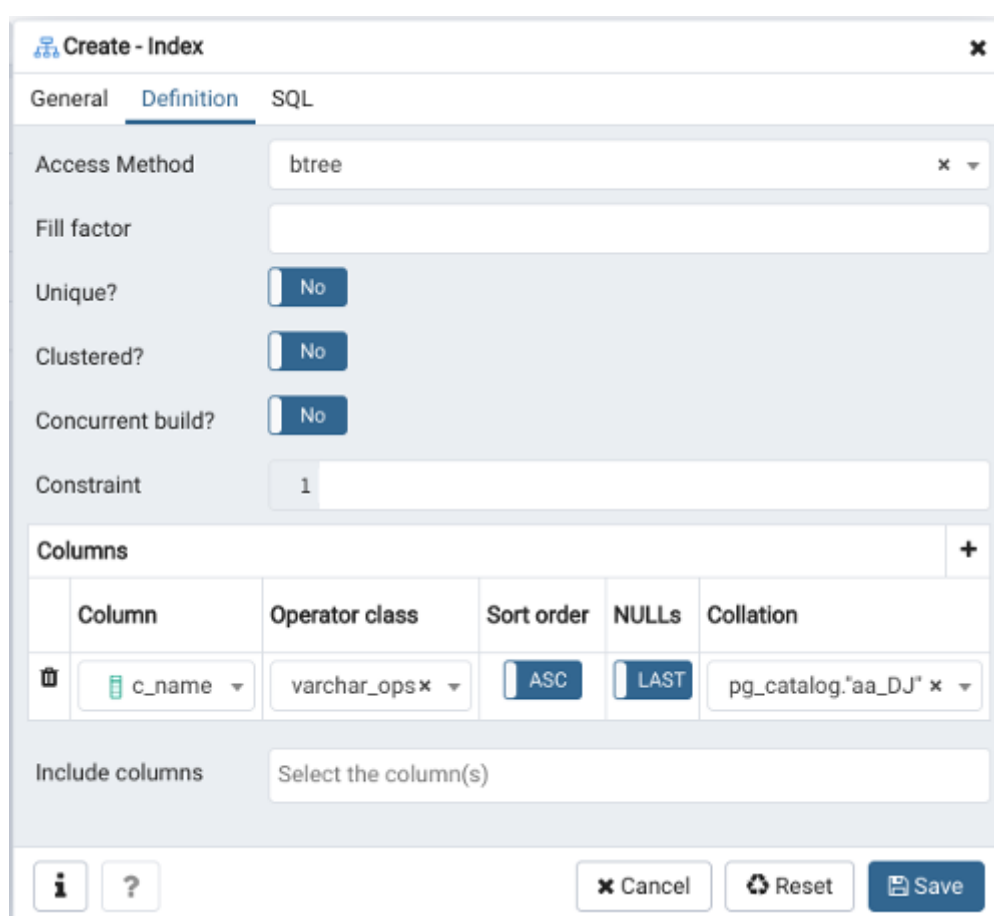
- Name:** A text field containing 'index_prod'.
- Tablespace:** A drop-down listbox showing 'pg_default'.
- Comment:** A large text area for notes.

At the bottom, there are buttons for 'Cancel', 'Reset', and 'Save', along with information and help icons.

Use the fields in the *General* tab to identify the index:

- Use the *Name* field to add a descriptive name for the index. The name will be displayed in the *pgAdmin* tree control.
- Use the drop-down listbox next to *Tablespace* to select the tablespace in which the index will reside.
- Store notes about the index in the *Comment* field.

Click the *Definition* tab to continue.



The 'Create - Index' dialog box is shown with the 'Definition' tab selected. It contains the following fields:

- Access Method:** A drop-down listbox showing 'btree'.
- Fill factor:** A text field.
- Unique?:** A toggle button set to 'No'.
- Clustered?:** A toggle button set to 'No'.
- Concurrent build?:** A toggle button set to 'No'.
- Constraint:** A text field containing '1'.
- Columns:** A table with columns: Column, Operator class, Sort order, NULLs, and Collation.

| | Column | Operator class | Sort order | NULLs | Collation |
|--|--------|----------------|------------|-------|--------------------|
| | c_name | varchar_ops | ASC | LAST | pg_catalog."aa_DJ" |
- Include columns:** A text field with the placeholder 'Select the column(s)'.

At the bottom, there are buttons for 'Cancel', 'Reset', and 'Save', along with information and help icons.

Use the fields in the *Definition* tab to define the index:

- Use the drop-down listbox next to *Access Method* to select an index type:

- Select *btree* to create a B-tree index. A B-tree index may improve performance when managing equality and range queries on data that can be sorted into some ordering (the default).
- Select *hash* to create a hash index. A hash index may improve performance when managing simple equality comparisons.
- Select *gist* to create a GiST index. A GiST index may improve performance when managing values with more than one key.
- Select *gin* to create a GIN index. A GIN index may improve performance when managing two-dimensional geometric data types and nearest-neighbor searches.
- Select *spgist* to create a space-partitioned GiST index. A SP-GiST index may improve performance when managing non-balanced data structures.
- Select *brin* to create a BRIN index. A BRIN index may improve performance when managing minimum and maximum values and ranges.

- Use the *Fill Factor* field to specify a fill factor for the index. The fill factor specifies how full the selected method will try to fill each index page.
- Move the *Unique?* switch to the *Yes* position to check for duplicate values in the table when the index is created and when data is added. The default is *No*.
- Move the *Clustered?* switch to the *Yes* position to instruct the server to cluster the table.
- Move the *Concurrent build?* switch to the *Yes* position to build the index without taking any locks that prevent concurrent inserts, updates, or deletes on the table.
- Use the *Constraint* field to provide a constraint expression; a constraint expression limits the entries in the index to those rows that satisfy the constraint.

Use the context-sensitive fields in the *Columns* panel to specify which column(s) the index queries. Click the *Add* icon (+) to add a column:

- Use the drop-down listbox in *Column* field to select the name of the column from the table.
- If enabled, use the drop-down listbox to select an available *Operator class* to specify the type of action performed on the column.
- If enabled, move the *Sort order* switch to specify the sort order:

- Select *ASC* to specify an ascending sort order (the default);
- Select *DESC* to specify a descending sort order.

- If enabled, move the *Nulls* switch to specify the sort order of nulls:

- Select *First* to specify nulls sort before non-nulls;
- Select *Last* to specify nulls sort after non-nulls (the default).

- Use the drop-down listbox in the *Collation* field to select a collation to use for the index.

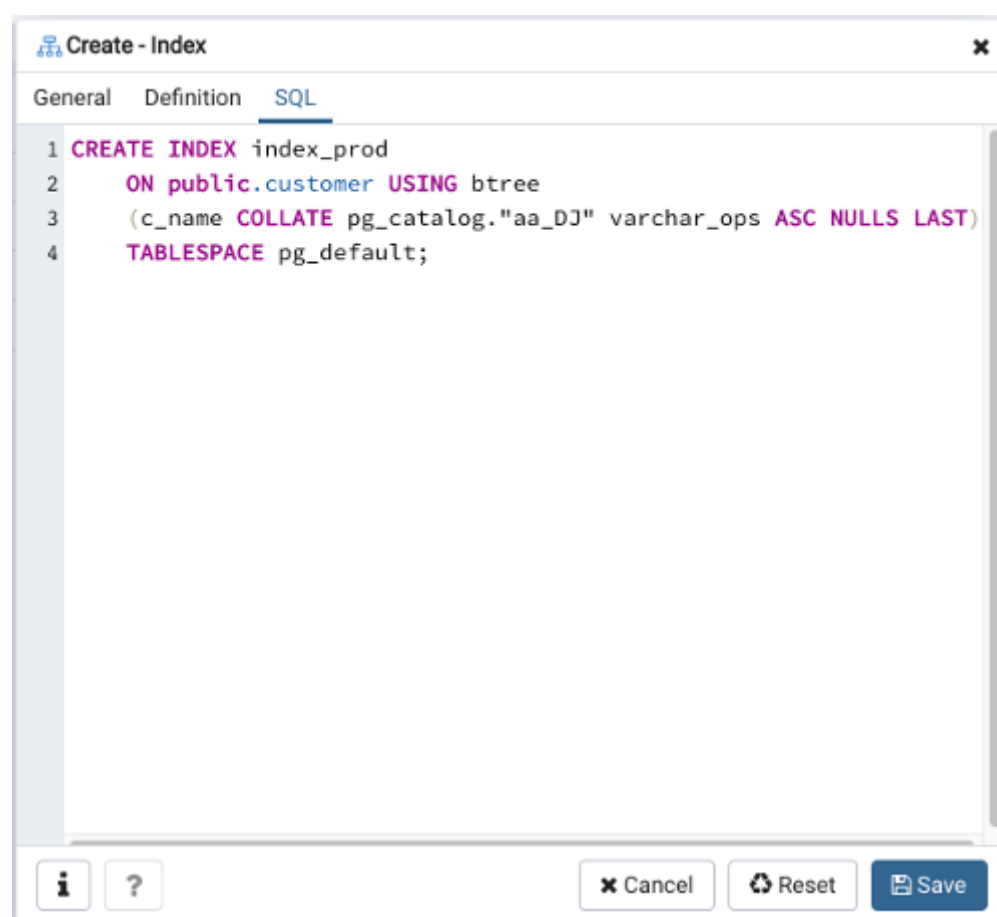
Use *Include columns* field to specify columns for *INCLUDE* clause of the index. This option is available in Postgres 11 and later.

Click the *SQL* tab to continue.

Your entries in the *Index* dialog generate a SQL command (see an example below). Use the *SQL* tab for review; revisit or switch tabs to make any changes to the SQL command.

Example ¶

The following is an example of the sql command generated by user selections in the *Index* dialog:



The example shown demonstrates creating an index named *dist_codes* that indexes the values in the *code* column of the *distributors* table.

- Click the *Info* button (i) to access online help.
- Click the *Save* button to save work.
- Click the *Cancel* button to exit without saving work.
- Click the *Reset* button to restore configuration parameters.